

# Cache

Computer Engineering 2

- **Principle of locality – The memory hierarchy**
- **Cache mechanics**
- **Cache organization**
  - Fully associative
  - Direct mapped
  - N-way set associative
- **Performance**
- **Replacement and write strategies**
- **The programmer's perspective**

## ■ Situation

- Processor
  - Fast cycle time
- Fast DRAM<sup>1</sup>
  - Slow cycle time (up to 100x)
  - Efficiently reads only in bursts
- → Bridging the gap such that pipelining is effective!

## ■ Goal

- Access “slower” memory in bursts and maintain a fast cache for fast single accesses
- But: Data consistency must be carefully managed, such that both, cache and main memory have the same data

<sup>1</sup> *Dynamic RAM*

- **At the end of this lesson you will be able**
  - to understand the principles of cache memory
  - to explain the principle of locality
  - to enumerate advantages and disadvantages of different cache models
    - Fully associative
    - Direct mapped
    - N-way set associative
  - to enumerate types of cache misses
  - to understand how cache size and cache hit rate are related
  - to name different replacement strategies

## ■ Principle of locality

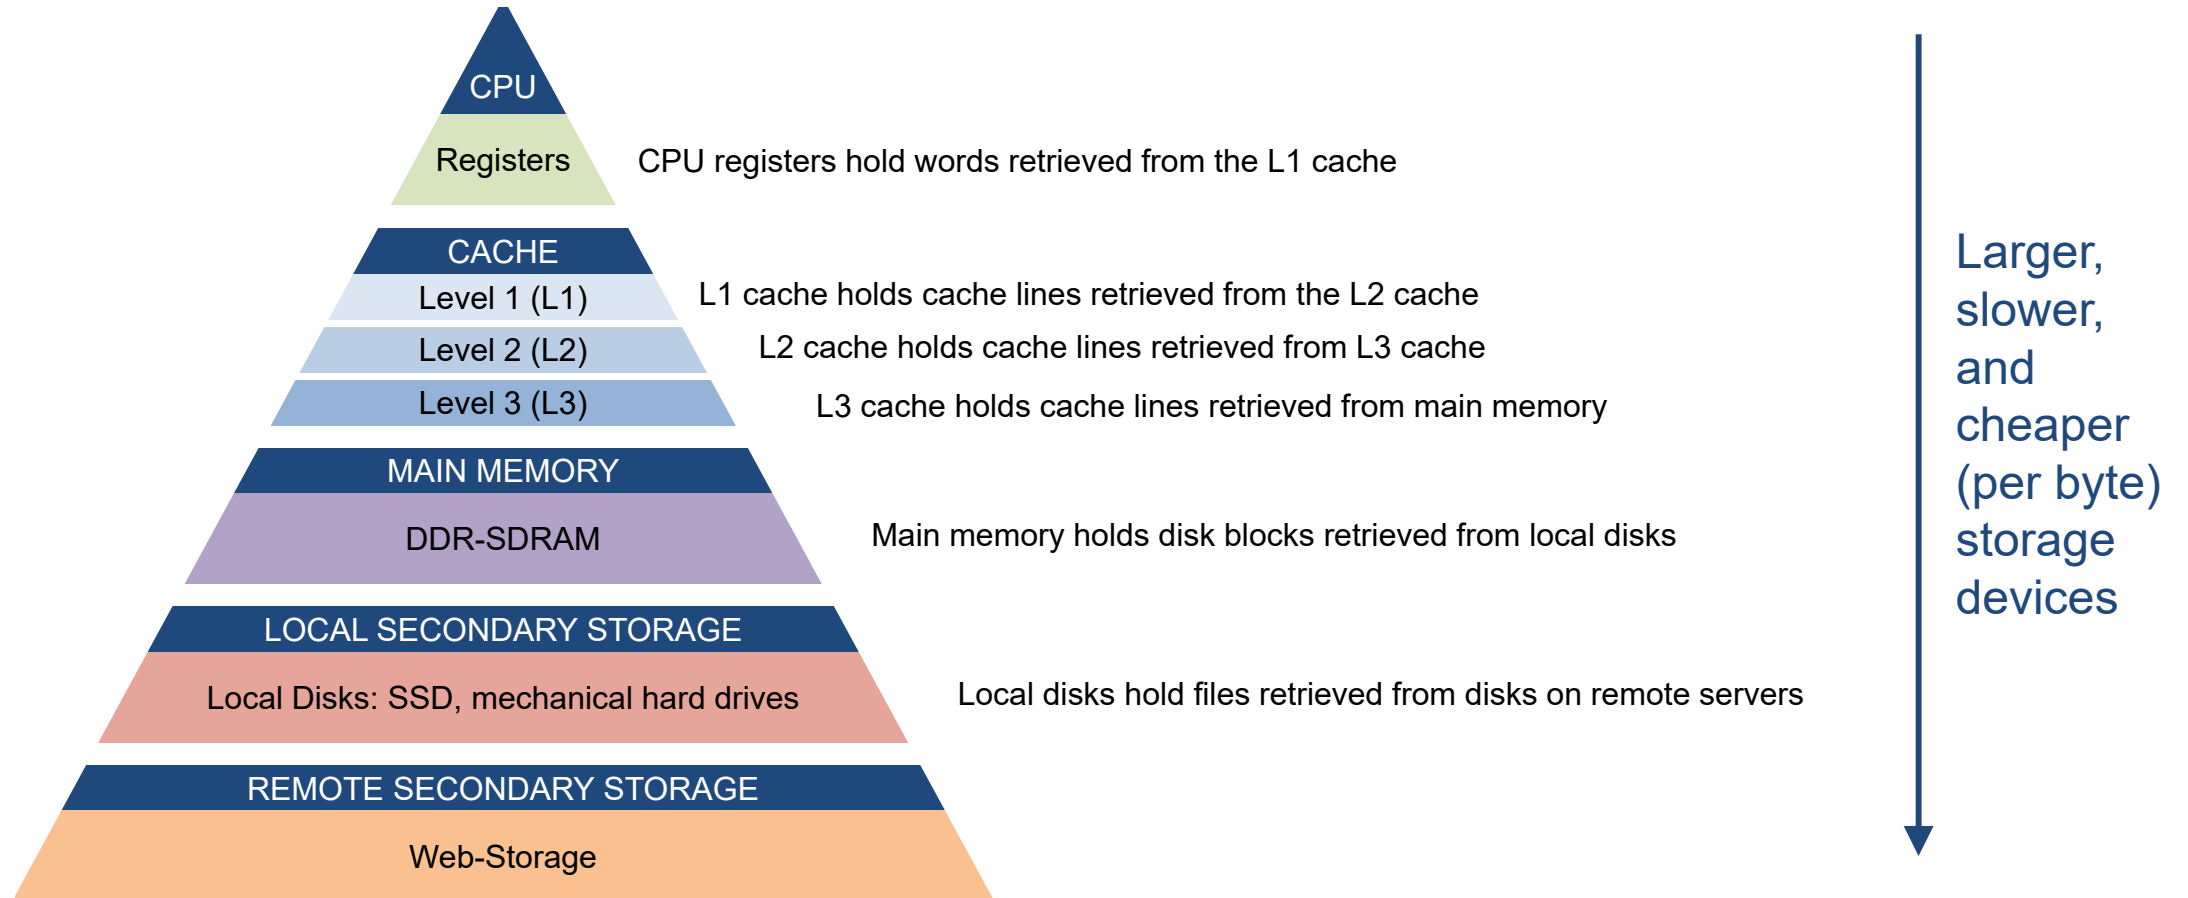
Programs usually access small regions of memory in a given interval of time

- **Spatial locality** → Current data location is likely being close to next accessed location
- **Temporal locality** → Current data location is likely being accessed again in near future

```
for(i = 0; i < 100000; i++) {           // incremental access
    a[i] = b[i];                         // → spatial locality
}

if (a[1234] == a[4321]) {                // → temporal locality
    a[1234] = 0;
}
```

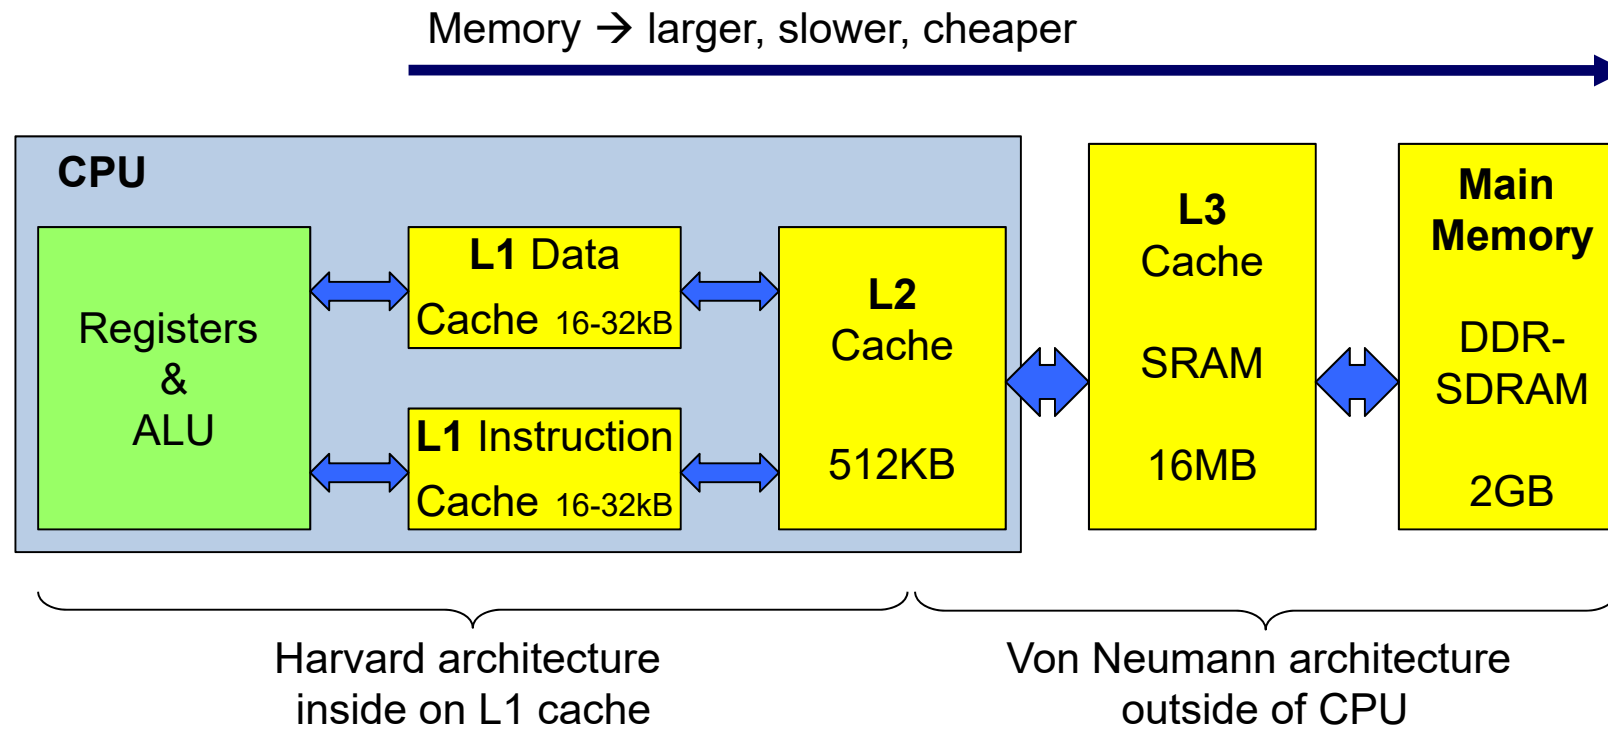
# The Memory Hierarchy



Adapted from Bryant and O'Hallaron

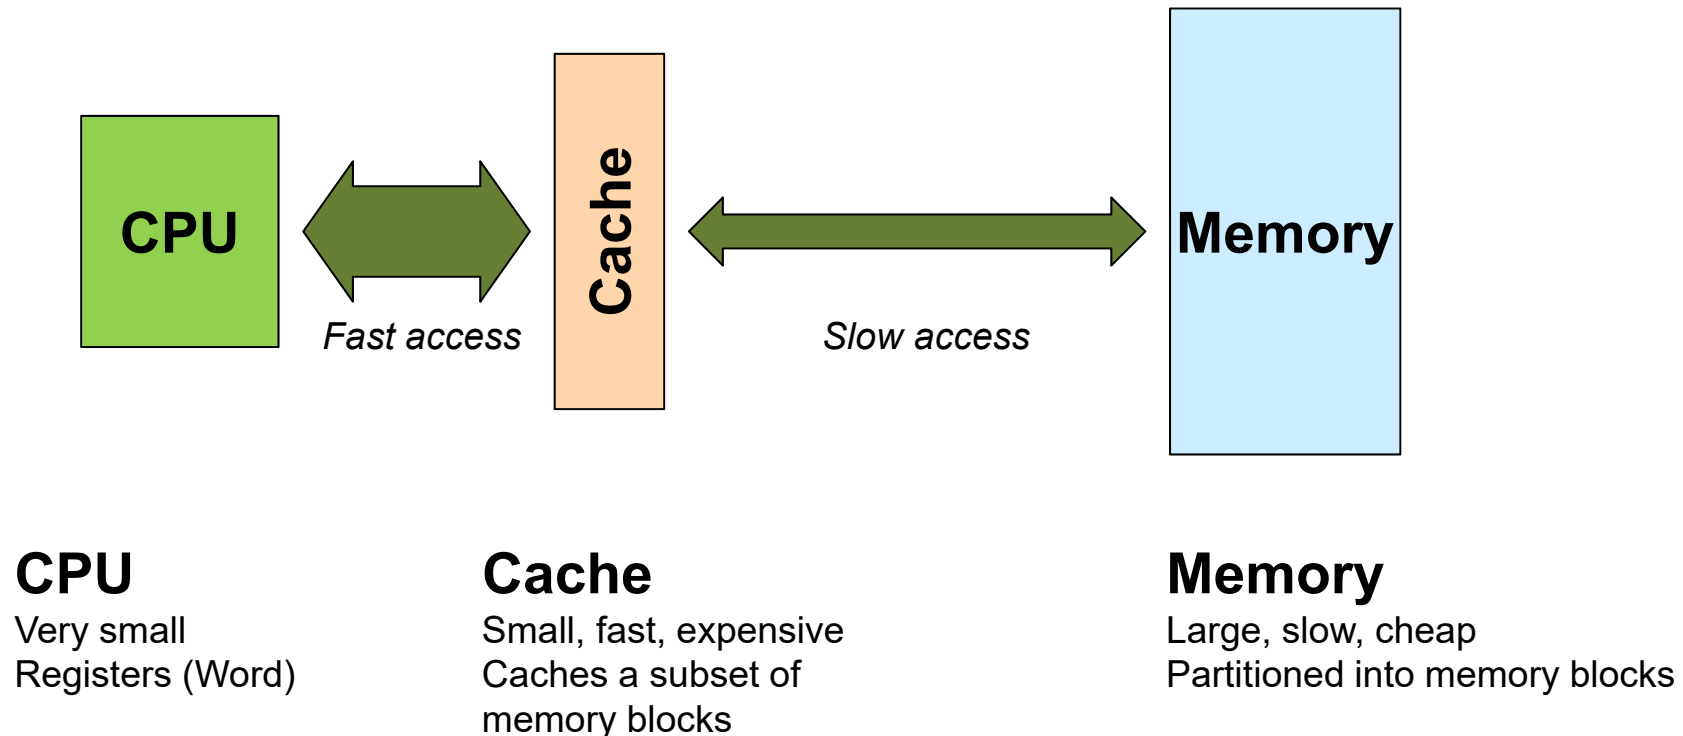
## ■ Cache Levels

- Typical Cache Architecture



## ■ Definition cache

- Computer memory with short access time
- Storage of frequently or recently used instructions or data



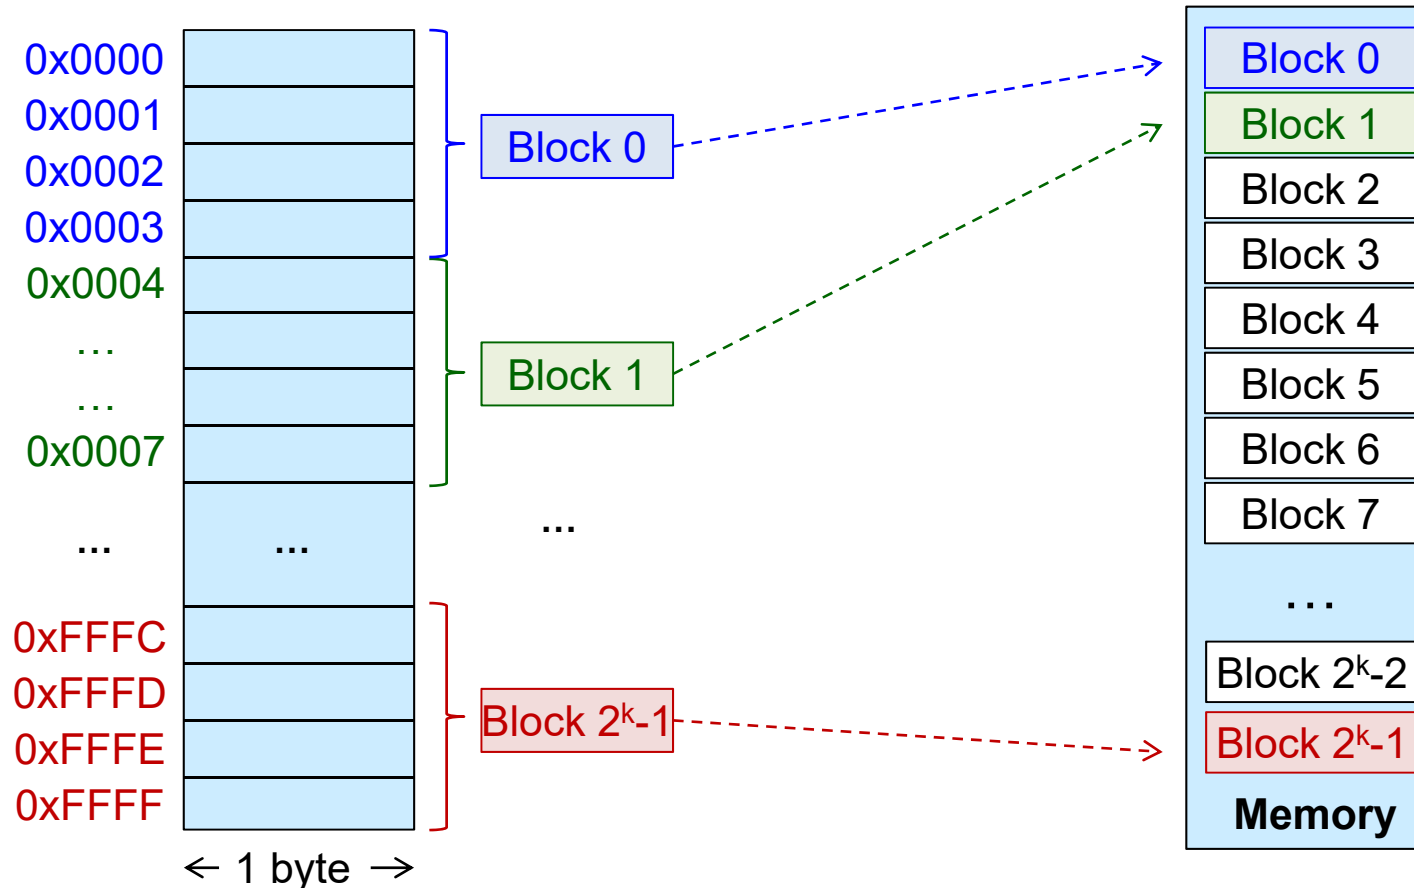


## ■ Memory blocks

- Address range is partitioned into memory blocks

*All examples given in this lecture are using*

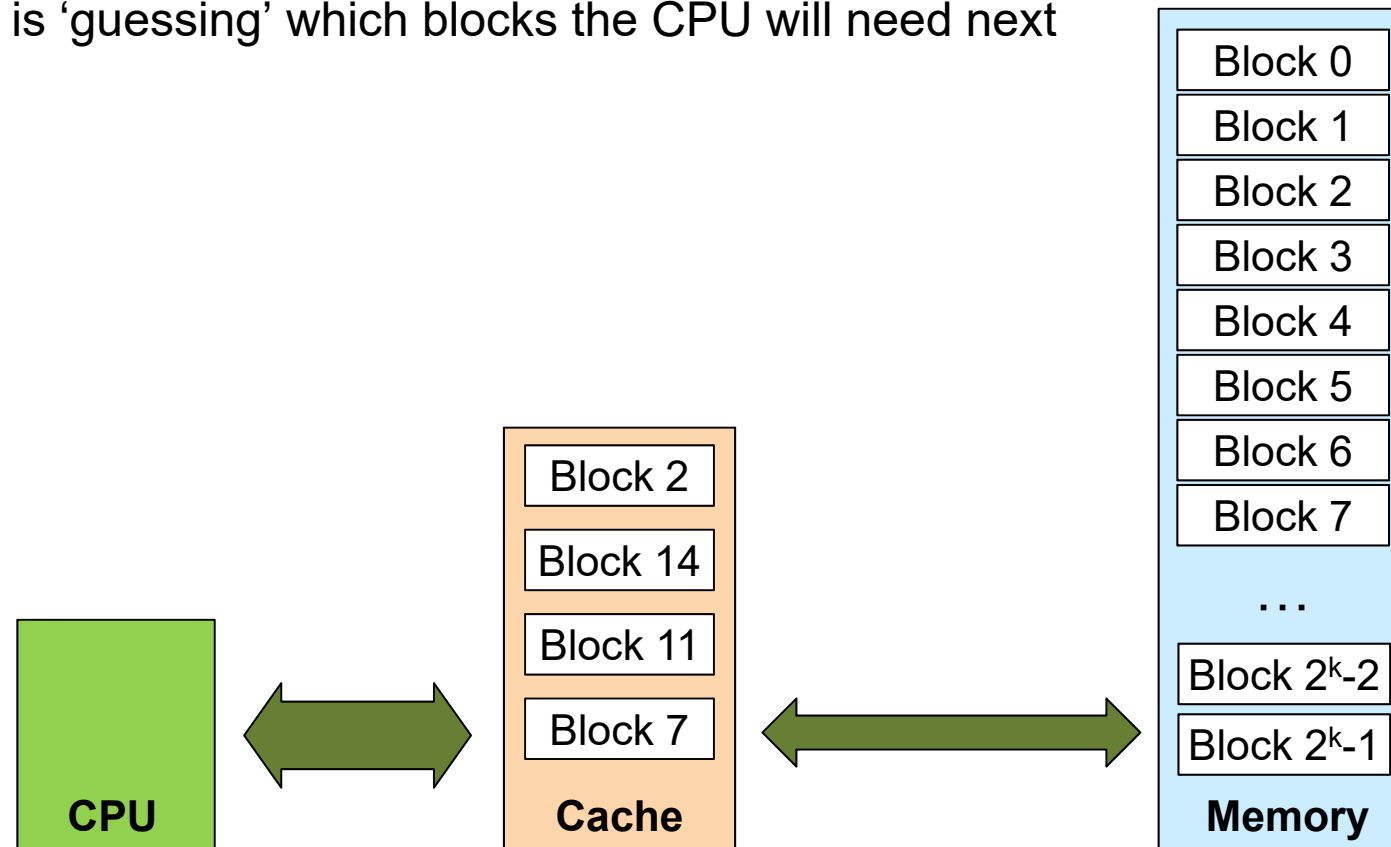
- *hypothetical 16 bit wide addresses*
- *blocks of 4 Bytes*



The number of Bytes per block is a design decision

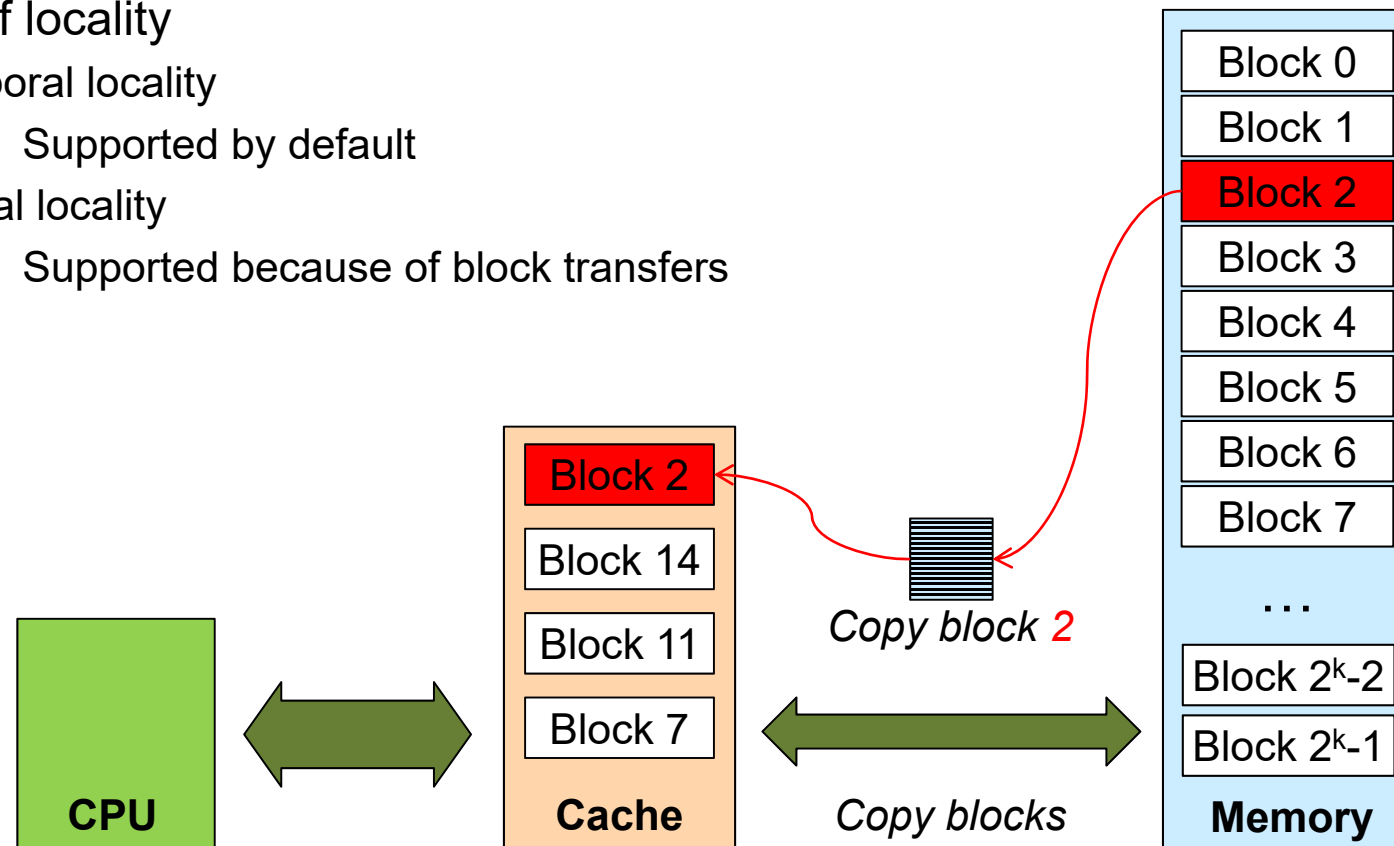
## ■ Memory blocks

- Selected blocks of main memory copied to faster cache memory
- The cache is 'guessing' which blocks the CPU will need next




## ■ Memory blocks

- Blocks of main memory copied to faster cache memory
- Principle of locality
  - Temporal locality
    - ▶ Supported by default
  - Spatial locality
    - ▶ Supported because of block transfers




## ■ Cache hit

- Data in block 14 is needed
- Block 14 in cache → **Hit**

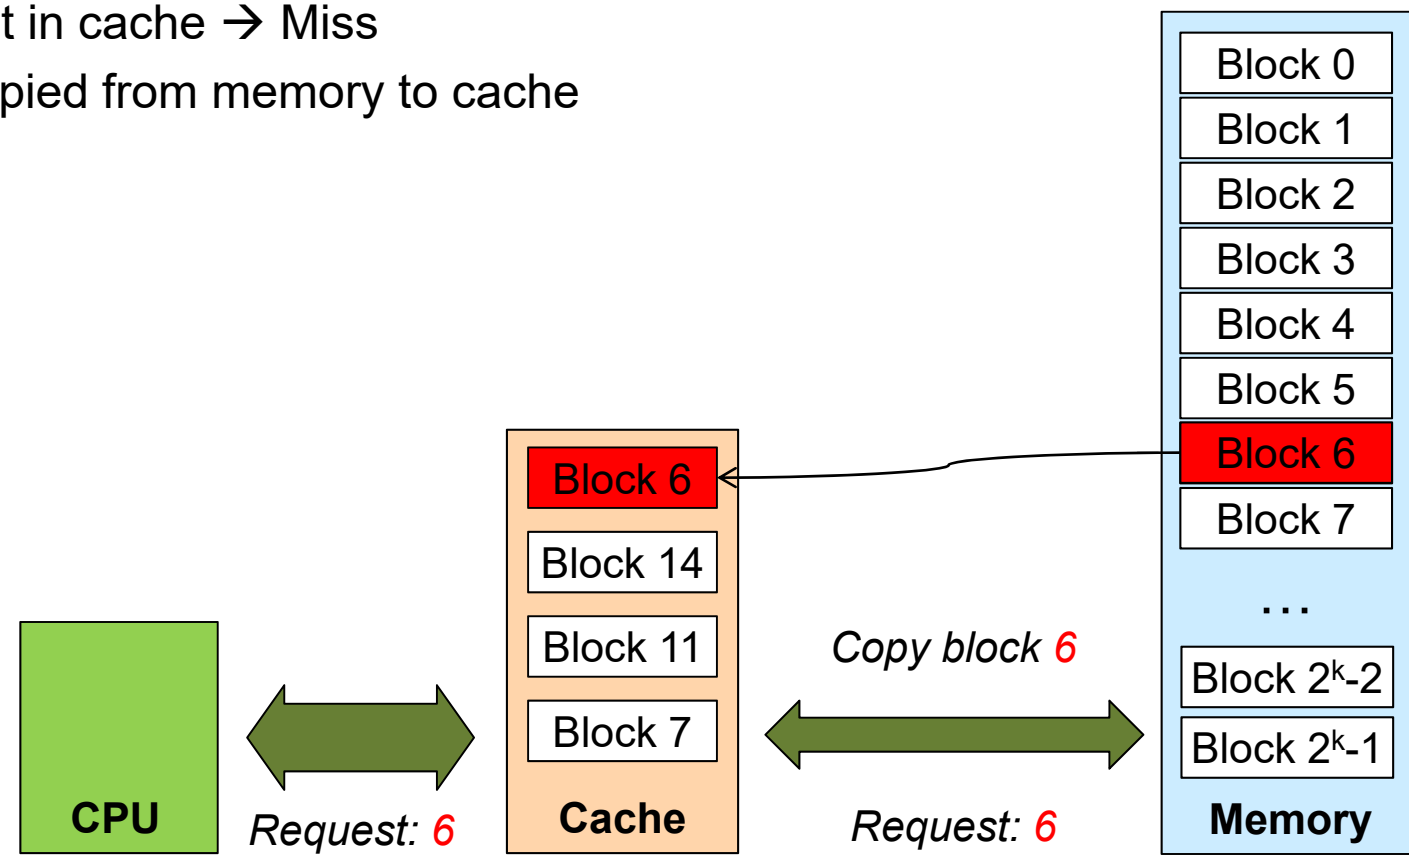


- **Cache miss (I)**
  - Data in block 6 is needed
  - Block 6 not in cache → **Miss**



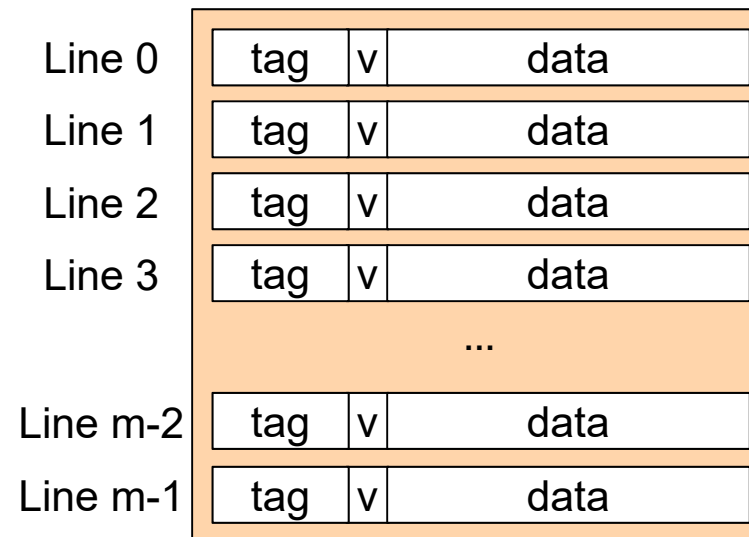
## ■ Cache miss (II)

- Data in block 6 is needed
- Block 6 not in cache → Miss
- Block 6 copied from memory to cache




## ■ Organized in lines

- Valid bit  $v$  → indicates that line contains valid data
- Tag → unique identifier for memory location
- Data → data of exactly one memory block
- $m$  = overall number of cache lines



## ■ Addressing

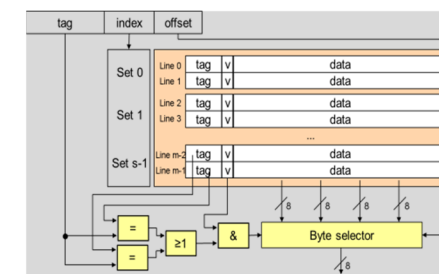
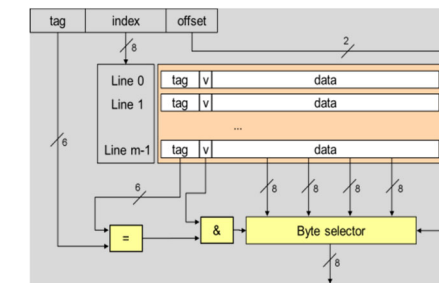
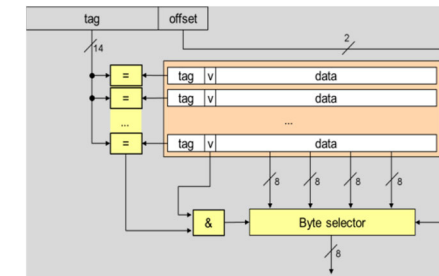


Block	Address	Block identification bits	offset
Block 0	0x0000	0000 0000 0000 00	00
	0x0001	0000 0000 0000 00	01
	0x0002	0000 0000 0000 00	10
	0x0003	0000 0000 0000 00	11
Block 1	0x0004	0000 0000 0000 01	00
	0x0005	0000 0000 0000 01	01
	0x0006	0000 0000 0000 01	10
	0x0007	0000 0000 0000 01	11
...			
Block 2k - 1	0xFFFC	1111 1111 1111 11	00
	0xFFFD	1111 1111 1111 11	01
	0xFFFE	1111 1111 1111 11	10
	0xFFFF	1111 1111 1111 11	11



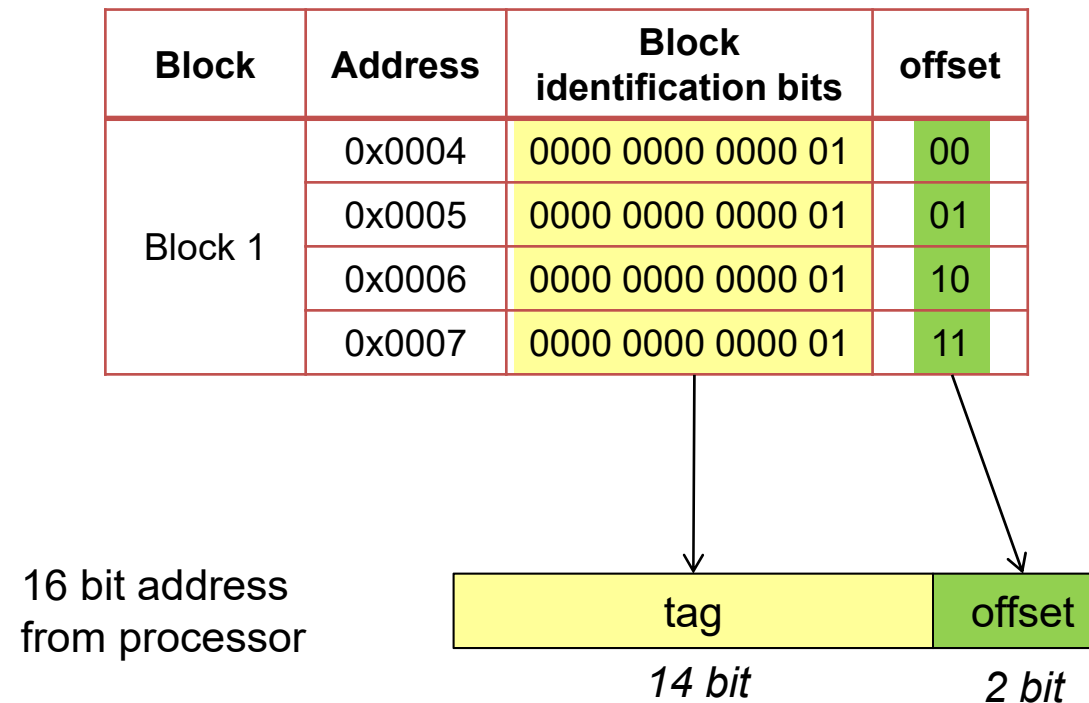
## ■ Three different cache models

- Fully associative
- Direct mapped
- N-way set associative



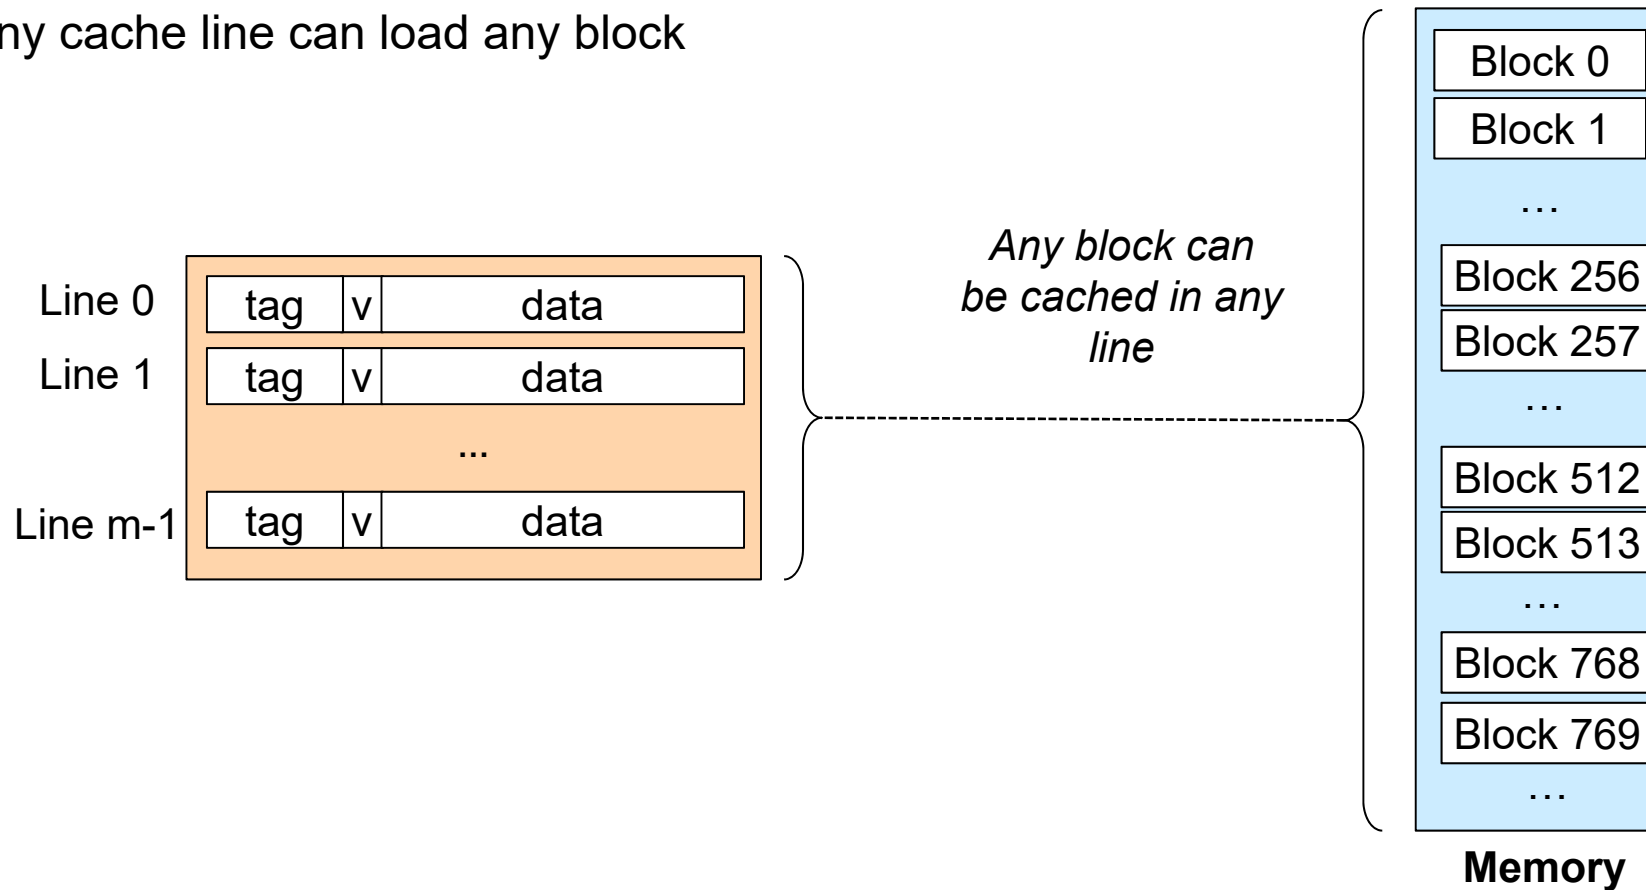
## ■ Addressing

- Tag contains complete block identification

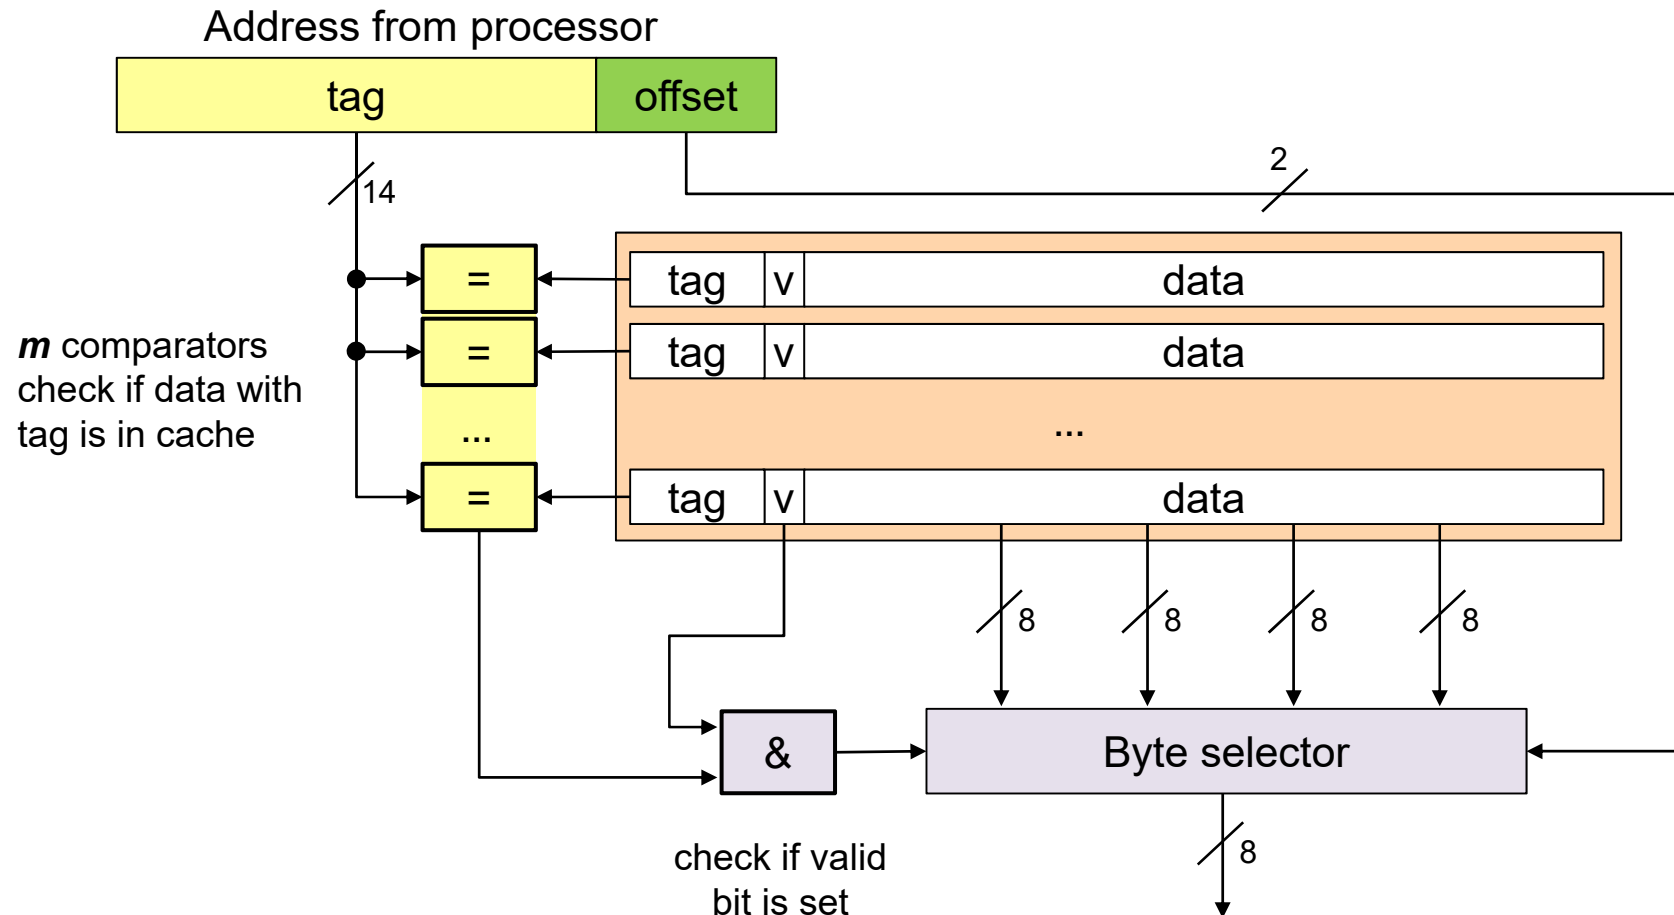


## ■ Organization

- Tag contains complete block identification
- Any cache line can load any block

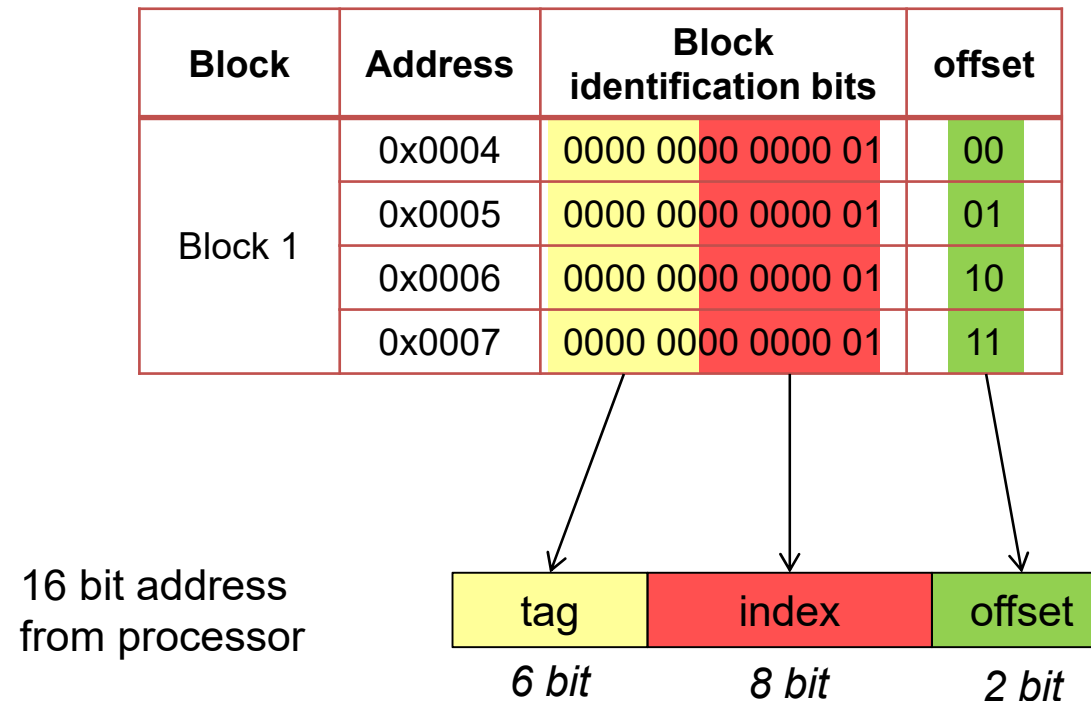


## ■ Architecture



## ■ Addressing

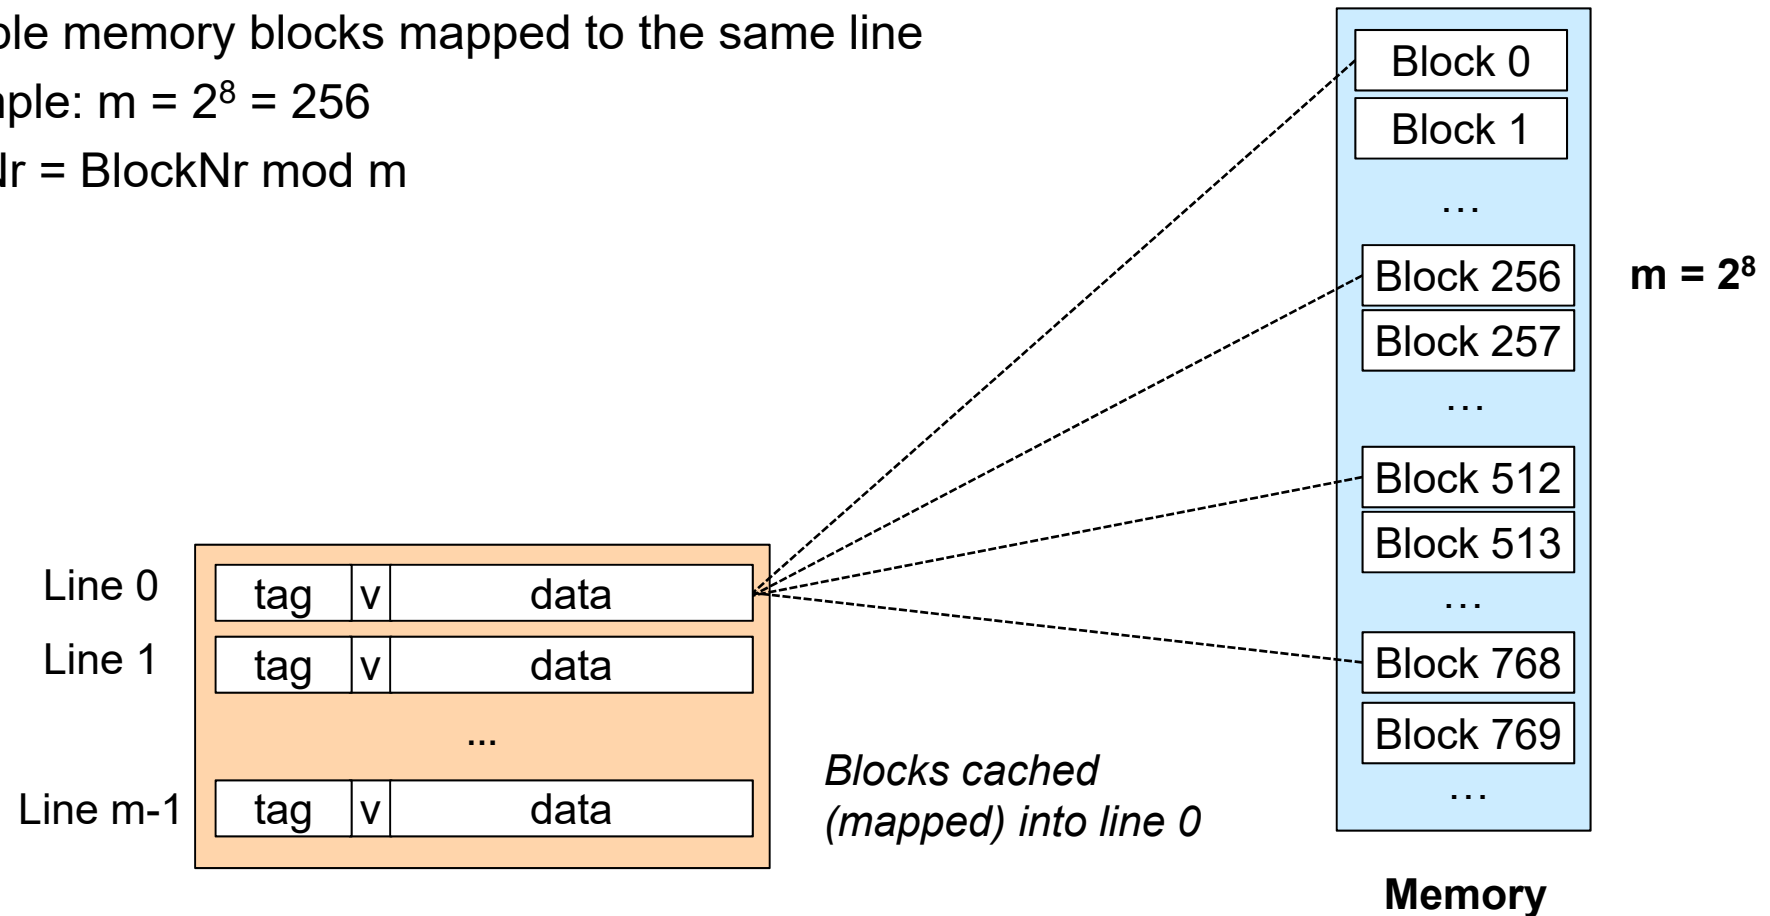
- Block identification split into tag and index



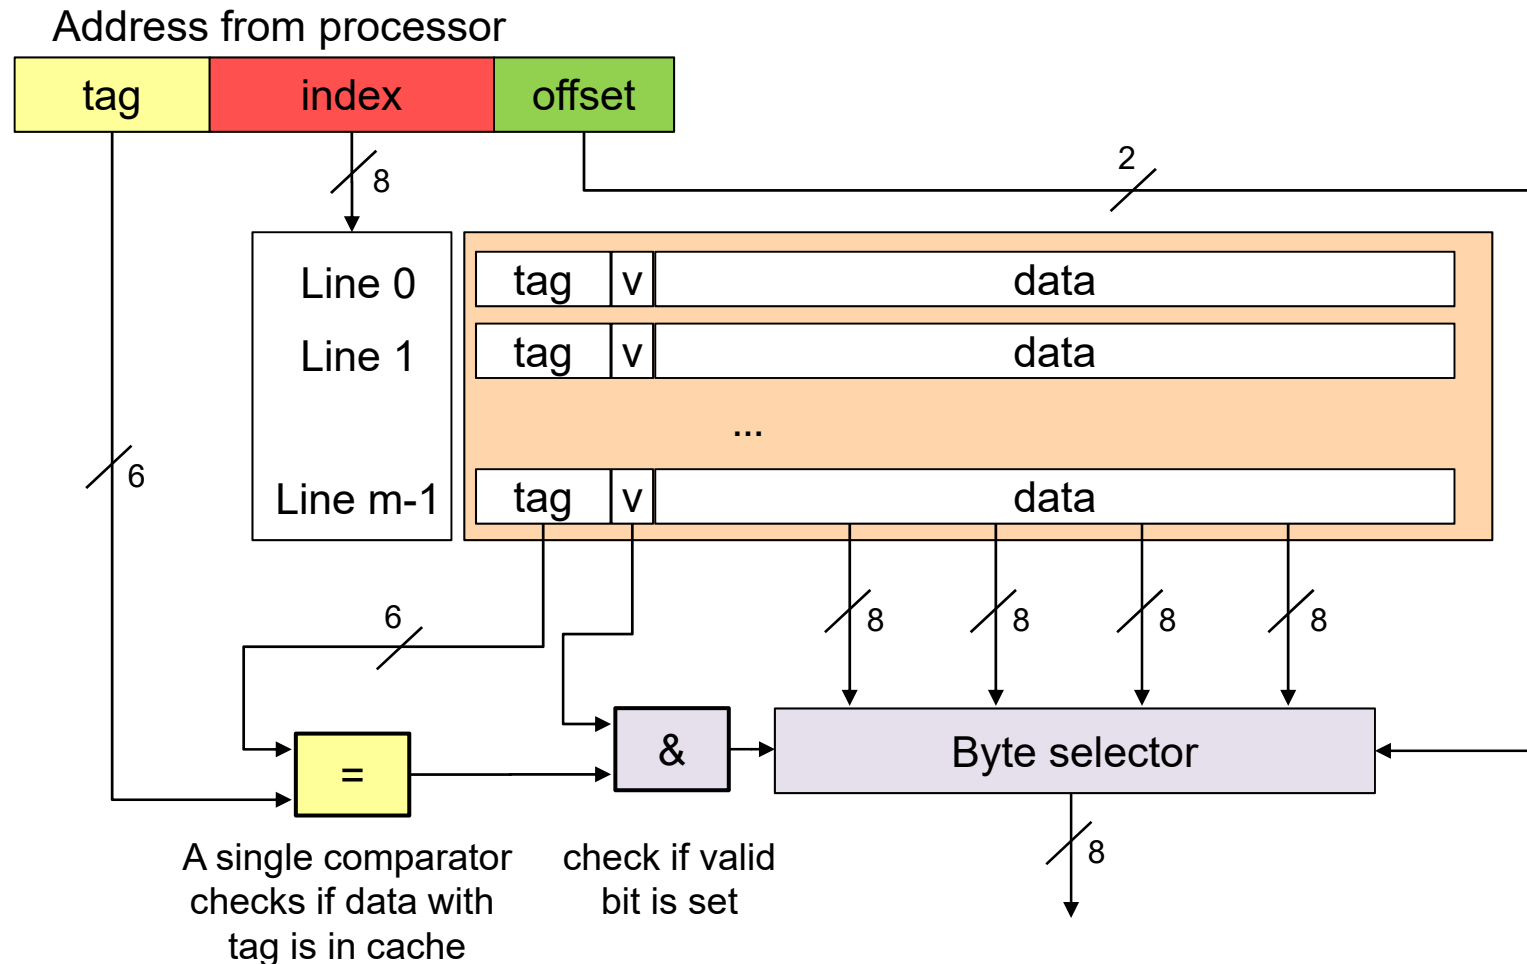
Bit sizes are an example

## ■ Organization

- Each memory block is mapped to exactly one cache line
- Multiple memory blocks mapped to the same line
- Example:  $m = 2^8 = 256$
- $\text{LineNr} = \text{BlockNr} \bmod m$



## ■ Architecture



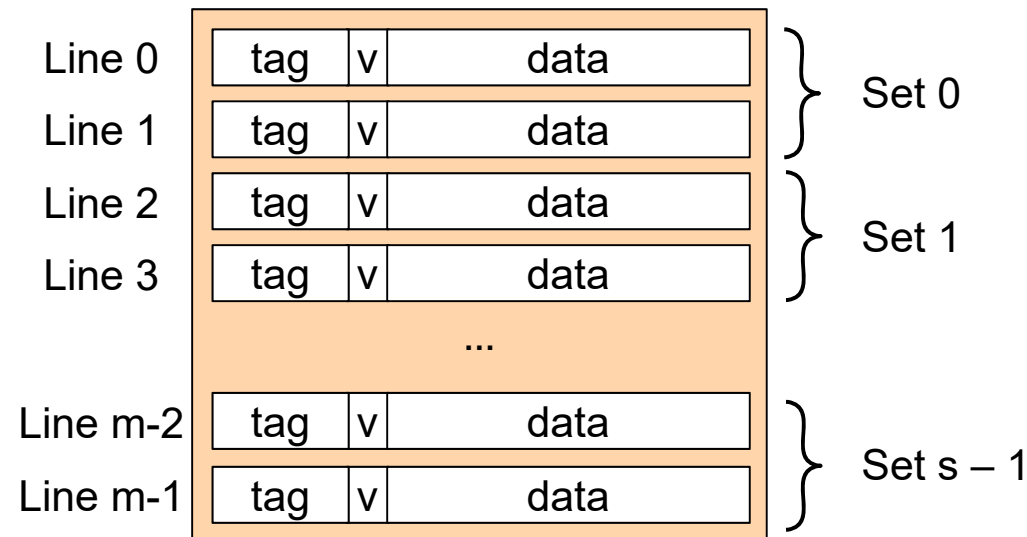
## ■ Organization

- Partition into sets
  - $s = m/n$  number of sets
  - $n$  lines per set („N-way“)
  - $b$  Bytes per line
- $s \times n \times b$  data Bytes

### N-Way Set Associative

Finding a compromise between simple logic (Direct Mapped) and high hit rates (Fully Associative)

*Example:  $k = 2$*






## ■ Addressing

- Example:  $n = 2$ 
  - Maximum index corresponds to number of sets ( $s = m/n$ )

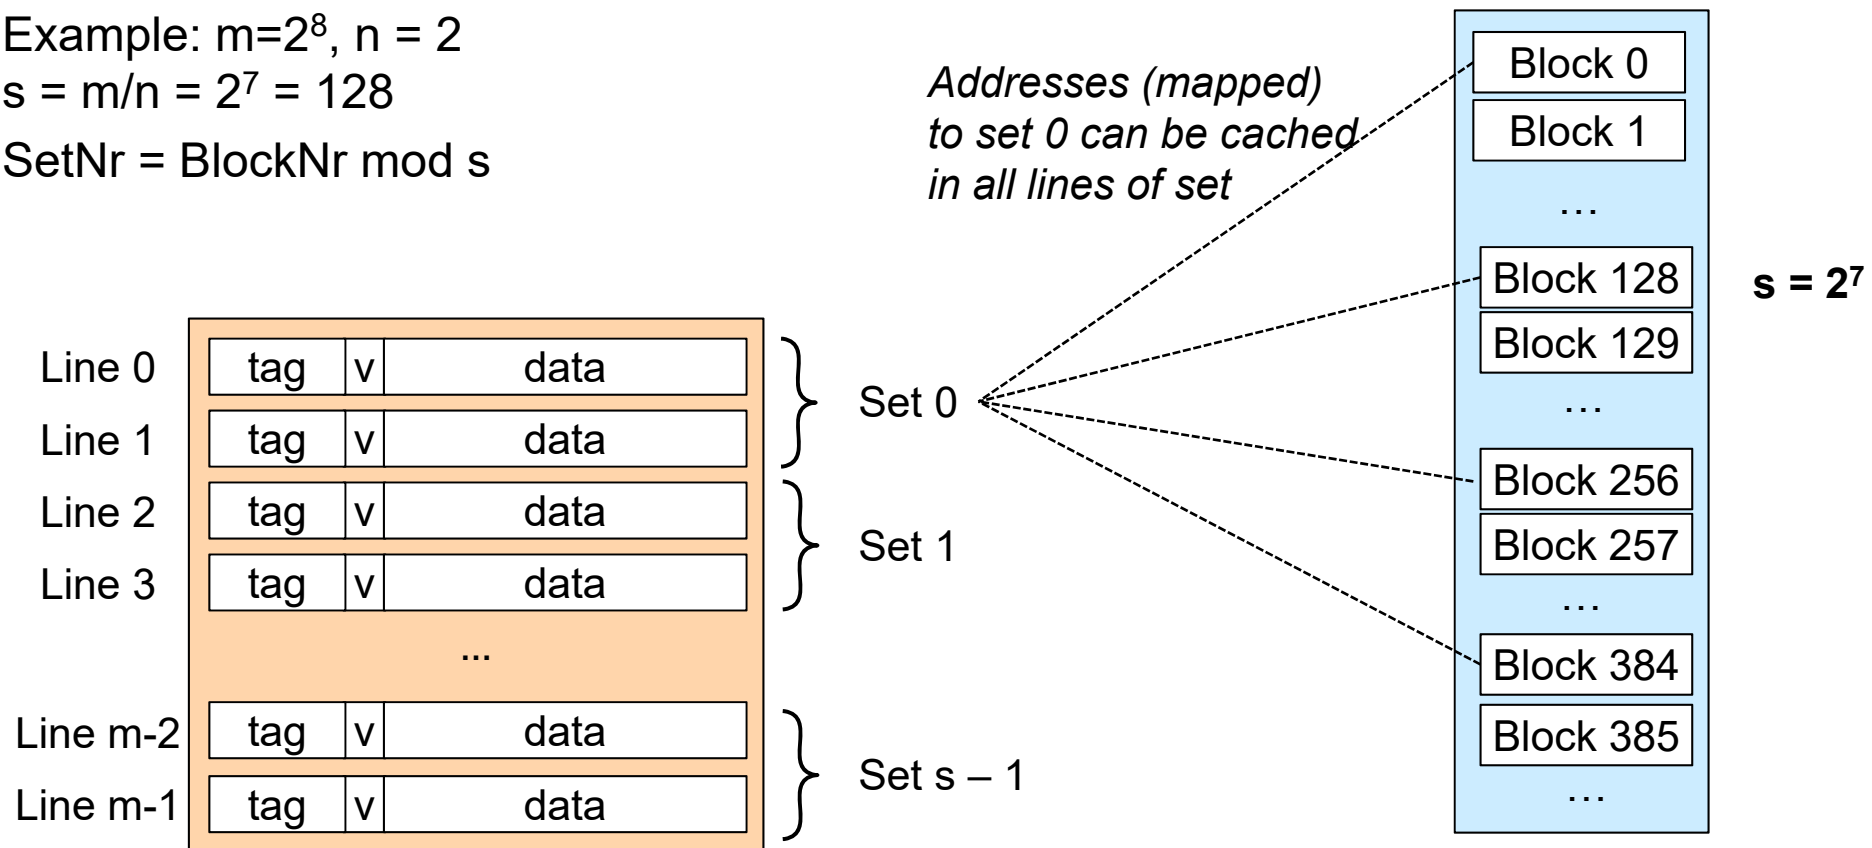
Block	Address	Block identification bits	offset
Block 1	0x0007	0000 0000 0000 01	11
	0x0006	0000 0000 0000 01	10
	0x0005	0000 0000 0000 01	01
	0x0004	0000 0000 0000 01	00

16 bit address  
from processor



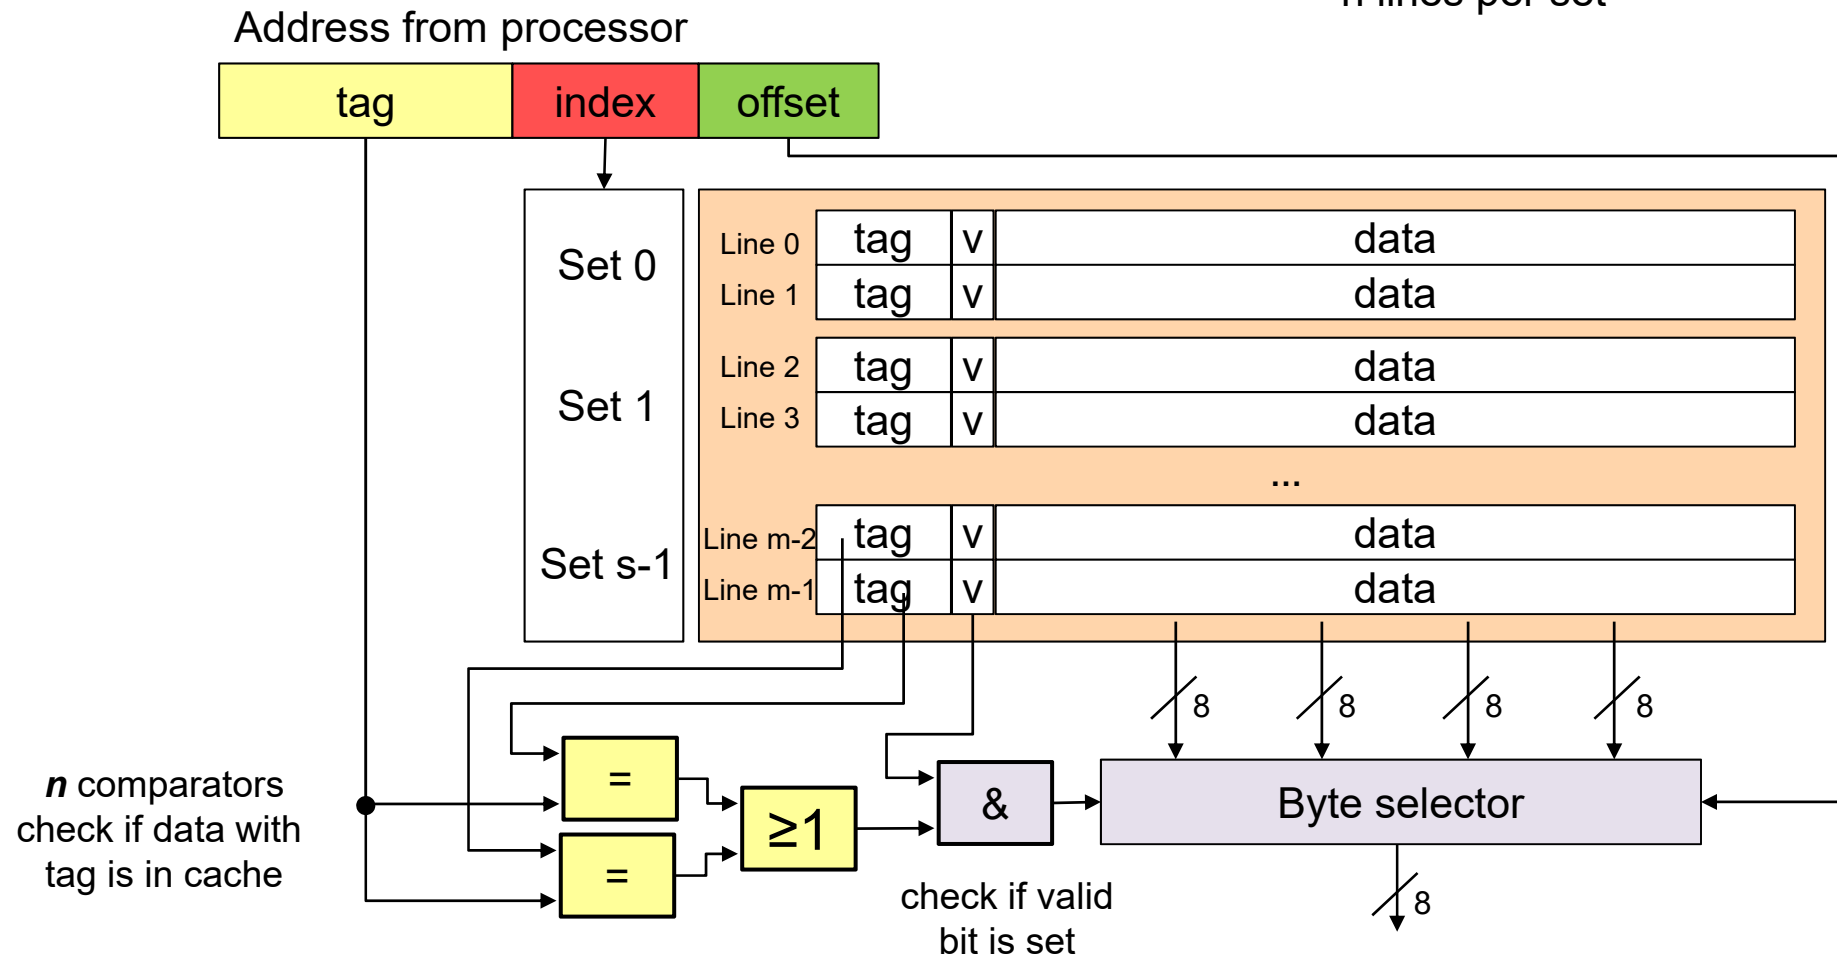
## ■ Organization

- Partition into sets
- Example:  $m=2^8$ ,  $n = 2$   
 $s = m/n = 2^7 = 128$
- $\text{SetNr} = \text{BlockNr} \bmod s$



## ■ Architecture

Example:  $n = 2$   
 $s = m/n$  number of sets  
 $n$  lines per set



## ■ Comparison

Organization	Fully associative	Direct mapped	N-way set associative
Number of sets	1	m	m/n
Associativity	m(=n)	1	n
Advantages	<ul style="list-style-type: none"><li>• Fast, flexible</li><li>• Highest hit rates</li><li>• Advanced replacement strategies</li></ul>	<ul style="list-style-type: none"><li>• Simple logic</li><li>• Replacement strategy defined by organization</li></ul>	Combination of both other concepts to combine advantages and to compensate disadvantages
Disadvantages	<ul style="list-style-type: none"><li>• Complex logic: one comparator per line</li><li>• Requires large area on silicon</li><li>• Replacement can be complex</li></ul>	<ul style="list-style-type: none"><li>• Lower hit rates</li></ul>	

# Performance – Cache Misses

---

- **Cold miss**
  - First access to a block
  
- **Capacity miss**
  - Working set larger than cache
  
- **Conflict miss**
  - Multiple data objects map to same slot

## ■ Hit rate / miss rate

- Fraction of memory references found / not found in cache
  - hit rate =  $\text{nr\_of\_hits} / \text{nr\_of\_accesses}$
  - miss rate =  $\text{nr\_of\_misses} / \text{nr\_of\_accesses} = 1 - \text{hit rate}$

## ■ Hit time

- Time to deliver a block in the cache to the processor

## ■ Miss penalty

- Additional time required to fetch data from memory because of a cache miss

## ■ High cache hit rate is important!


- 99% hit rate can be twice as fast as 97%

## ■ Example

- Cache hit time: 1 processor cycle
- Miss penalty: 100 processor cycles
- Average access time is:
  - 97% hits:  $1 \text{ cycle} + 0.03 * 100 \text{ cycles} = 4 \text{ cycles average}$
  - 99% hits:  $1 \text{ cycle} + 0.01 * 100 \text{ cycles} = 2 \text{ cycles average}$

## ■ Cache size versus hit rate

- Typical hit rate for Associativity 1, 2, 4, 8, 16 and cache size





## ■ Selecting cache line to replace by

- LRU: Least recently used
- LFU: Least frequently used
- FIFO: First In–First Out → oldest
- Random Replace: randomly chosen

*additional information  
needed in cache*

*simple, but still  
good performance*

→ Only relevant for Fully- / N-way associative caches

→ Hard coded in cache implementation (=hardware)

## ■ What to do on a write hit\*?

- Write-through
  - Write immediately to memory
- Write-back
  - Delay write to memory until replacement of line (needs a valid bit)

## ■ What to do on a write miss\*\*?

- Write-allocate
  - Load line into cache (from memory) and update line in cache
- No-write-allocate
  - Writes immediately to memory

\* *Write hit: data already in cache*

\*\* *Write miss: data not in cache*

# The Programmer's Perspective


- **For loops over multi-dimensional arrays**
  - Example: matrices (2-dim arrays)
- **Change order of iteration to match layout**
  - Gets better spatial locality
  - Layout in C: last index changes first!

```
for(j = 0; j < 10000; j++){
    for(i = 0; i < 40000; i++){
        c[i][j] = a[i][j] + b[i][j];
    }
}
```

// a[i][j] and a[i+1][j]  
// are 10'000 elements apart

```
for(i = 0; i < 40000; i++){
    for(j = 0; j < 10000; j++){
        c[i][j] = a[i][j] + b[i][j];
    }
}
```

// a[i][j] and a[i][j+1]  
// are next to each other



# The Programmer's Perspective

## ■ Change order of iteration to match layout

```
j is inner
1047527424-1048575999 done

real    0m3.460s
user    0m3.452s
sys     0m0.000s

j is outer
1047527424-1048575999 done

real    0m18.509s
user    0m18.472s
sys     0m0.000s
```

```
#include<stdio.h>
#include<time.h>

int main( int argc, char** argv )
{
    int iterations = 1000;
    int size = 1024;

    int array[size][size];
    int v = 0;

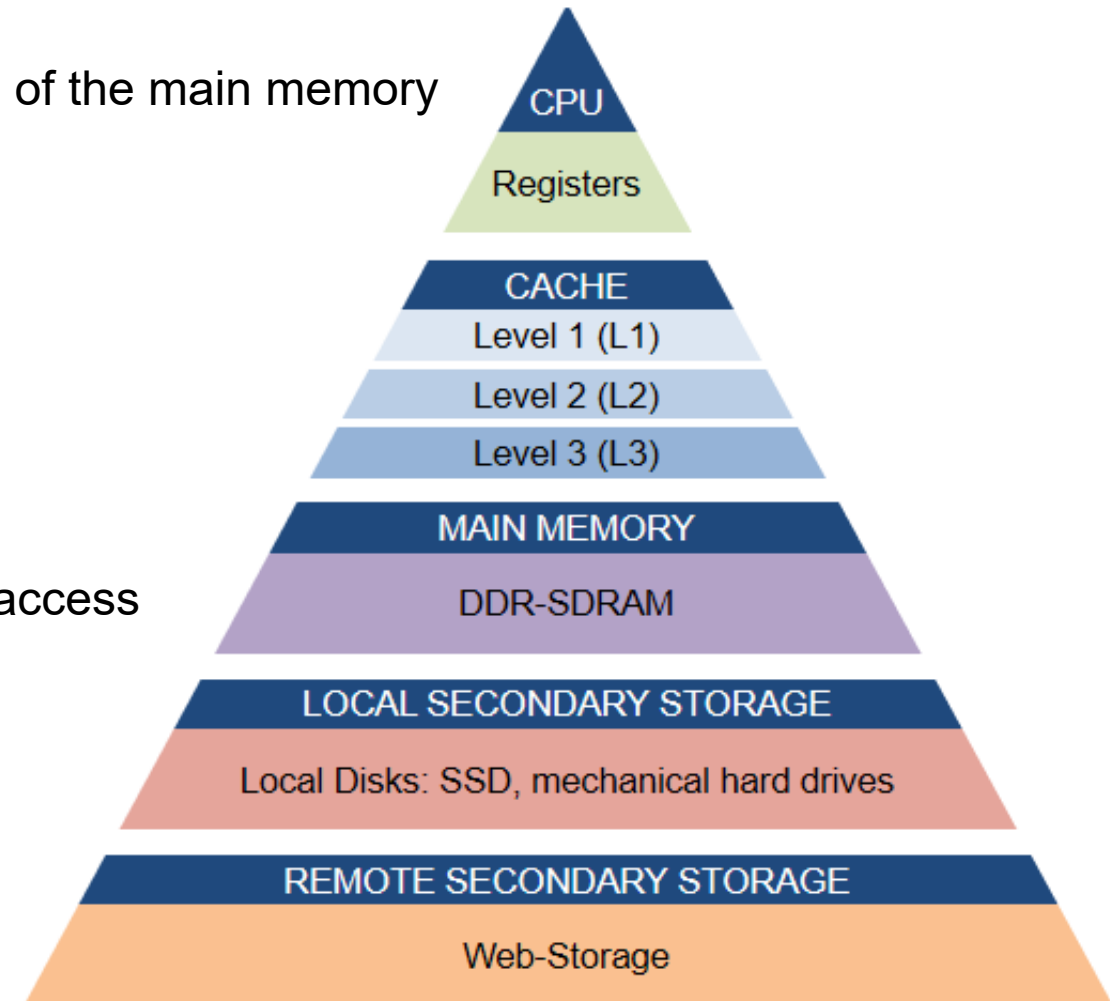
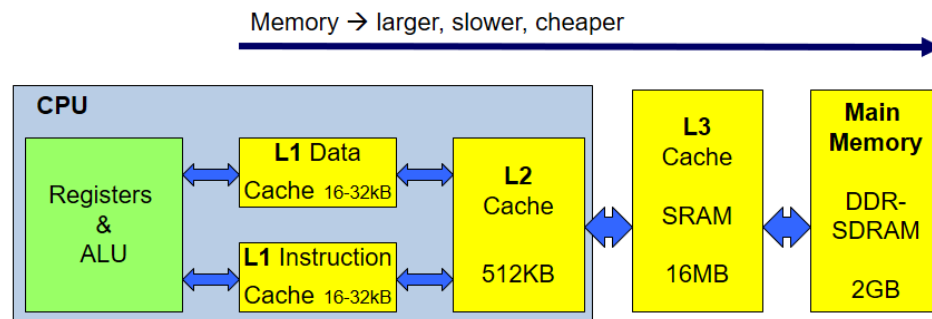
    if (argc==2) {
        printf("j is inner\n");
        for (int n=0; n<iterations; n++) {
            for (int i=0; i<size; i++) {
                for (int j=0; j<size; j++) {
                    array[i][j] = v;
                    v++;
                }
            }
        }
    }
    else {
        printf("j is outer\n");
        for (int n=0; n<iterations; n++) {
            for (int j=0; j<size; j++) {
                for (int i=0; i<size; i++) {
                    array[i][j] = v;
                    v++;
                }
            }
        }
    }

    printf("%i-%i done\n", array[0][0], array[size-1][size-1]);

    return 0;
}
```

## ■ Cache

- Cache allows fast data access to selected parts of the main memory which are mirrored in the cache
  - Spatial and temporal locality
- Different cache models
  - Fully associative
  - Direct mapped
  - N-way associative
- Replacement / Write strategies
- Reducing cache misses by optimizing memory access




# Driver for Cache and Pipelining

ACAT2013

IOP Publishing

Journal of Physics: Conference Series **523** (2014) 012002

doi:10.1088/1742-6596/523/1/012002



**Figure 1: CPU-memory performance gap. Modelled after "Computer Architecture": Hennessy, John L.; Patterson, David A.**

[https://www.researchgate.net/publication/273029990\\_Opportunities\\_and\\_choice\\_in\\_a\\_new\\_vector\\_era](https://www.researchgate.net/publication/273029990_Opportunities_and_choice_in_a_new_vector_era)


# Cache Architecture Comparison

	Characteristic	ARM Cortex-A8	Intel Nehalem
L1	L1 cache organization	Split instruction and data caches	Split instruction and data caches
	L1 cache size	32 KiB each for instructions/data	32 KiB each for instructions/data per core
	L1 cache associativity	4-way (I), 4-way (D) set associative	4-way (I), 8-way (D) set associative
	L1 replacement	Random	Approximated LRU
	L1 block size	64 bytes	64 bytes
	L1 write policy	Write-back, Write-allocate(?)	Write-back, No-write-allocate
	L1 hit time (load-use)	1 clock cycle	4 clock cycles, pipelined
L2	L2 cache organization	Unified (instruction and data)	Unified (instruction and data) per core
	L2 cache size	128 KiB to 1 MiB	256 KiB (0.25 MiB)
	L2 cache associativity	8-way set associative	8-way set associative
	L2 replacement	Random(?)	Approximated LRU
	L2 block size	64 bytes	64 bytes
	L2 write policy	Write-back, Write-allocate (?)	Write-back, Write-allocate
	L2 hit time	11 clock cycles	10 clock cycles
L3	L3 cache organization		Unified (instruction and data)
	L3 cache size		8 MiB, shared
	L3 cache associativity		16-way set associative
	L3 replacement		Approximated LRU
	L3 block size		64 bytes
	L3 write policy		Write-back, Write-allocate
	L3 hit time		35 Clock Cycles

Source: Patterson / Hennessey: Computer Organization and Design, 5th edition, Elsevier

# The CPU-Memory Gap

The gap widens between DRAM, disk, and CPU speeds.



Bryant and O'Hallaron, Computer Systems:  
A Programmer's Perspective, Third Edition